

cuRAMSES-kjhan

Technical Reference Manual

K-Section Ordering, Morton Key Octree,
Memory-Based Load Balancing, and
Performance Optimizations

Juhan Kim

February 2026

Based on RAMSES (R. Teyssier)

<git@github.com:kjhan0606/cuRAMSES-kjhan.git>

Contents

I Code Modifications

1	Quick Start Guide	9
1.1	Prerequisites	9
1.2	Building cuRAMSES	9
1.3	Setting Up a Zoom-In Simulation	9
1.3.1	Step 1: Generate Initial Conditions with MUSIC	9
1.3.2	Step 2: Create Zoom Geometry Scalar	10
1.3.3	Step 3: Create Run Directory	10
1.3.4	Step 4: Configure Namelist	11
1.3.5	Step 5: Run	11
1.4	Key Parameters Quick Reference	11
1.4.1	mass_cut_refine Reference	12
1.5	Troubleshooting	12
2	K-Section Ordering	15
2.1	Overview	15
2.2	Hierarchical MPI Exchange	15
2.2.1	Exclusive Exchange	15
2.2.2	Overlap Exchange	16
2.2.3	Periodic Boundary Conditions	16
2.3	Tree Navigation	16
2.4	Verification	17
3	Ghost Zone Exchange via K-Section	19
3.1	AMR Ghost Zones	19
3.1.1	Modified File	19
3.1.2	Subroutines	19
3.1.3	Data Packing	19
3.1.4	Dispatch	19
3.1.5	Bulk Exchange	20
3.2	build_comm via K-Section	20
3.3	MPI_ALLTOALL Replacements	20
3.4	Pre-Allocated Buffer Pool	20
4	Multigrid Poisson K-Section Communication	21
4.1	Motivation	21
4.2	Implementation	21
4.2.1	Key Differences from AMR Exchange	21

4.2.2	Forward Exchange	21
4.2.3	Reverse Exchange (Accumulation)	22
5	Morton Key Octree	23
5.1	Overview	23
5.1.1	Modified Files	23
5.2	Morton Key	23
5.3	Hash Table	23
5.4	Neighbor Lookup	23
5.5	nbor Array Removal (Phase 4)	23
6	Memory-Based Load Balancing	25
6.1	Motivation	25
6.2	Cost Function	25
6.3	Implementation Details	25
6.4	Parameters	25
7	Memory Savings: Large Array Optimization	27
7.1	Overview	27
8	IC Reading with Stream Access	29
8.1	Motivation	29
8.2	Implementation	29
8.2.1	Modified Files	29
9	Load Balance Profiling and Tuning	31
9.1	Internal Timing	31
9.2	nremap Tuning	31
9.3	Min/Max Memory Reporting	31
10	Zoom-In Simulation Setup	33
10.1	Overview	33
10.2	Initial Conditions	33
10.2.1	MUSIC Configuration	33
10.2.2	Generated IC Structure	34
10.3	Zoom Geometry Scalar (ic_pvar_00006)	34
10.3.1	The Problem	34
10.3.2	The Solution: Passive Scalar Mask	34
10.3.3	How It Works at Each Stage	35
10.3.4	Creating ic_pvar_00006	35
10.3.5	mass_cut_refine	36
10.4	Memory Considerations	36
10.4.1	ngridtot Sizing	36
10.4.2	CommitLimit	36

10.4.3	Virtual vs. Physical Memory	37
11	HDF5 I/O	39
11.1	Overview	39
11.2	Building HDF5	39
11.3	Namelist Parameters	40
11.4	HDF5 File Structure	40
11.4.1	Parallel Write Strategy	41
11.5	Implementation Files	41
11.5.1	Dispatch	42
11.6	Verification	42
11.7	Variable-NCPU Restart: Distributed Grid Creation	42
11.7.1	Distributed Algorithm	42
11.7.2	Key Design Decisions	43
11.7.3	Verification	43
11.8	Poisson MG Fine-Level Optimization	44
11.8.1	Precomputed Neighbor Grid Array	44
11.8.2	Merged Red-Black Gauss-Seidel Exchange	44
11.8.3	Residual and Norm Single Pass	44
11.8.4	Division to Multiplication	44
11.8.5	Performance Impact	45
12	Scaling Performance	47
12.1	Pure MPI Scaling	47
12.2	Hybrid MPI + OpenMP Scaling	47
12.3	Key Observations	48
12.4	Per-Routine Scaling Analysis	48
12.4.1	Scaling Categories	48
12.4.2	Bottleneck Shift	50
13	GPU Hydro Acceleration	51
13.1	Architecture Overview	51
13.2	Hybrid CPU/GPU Dispatch	51
13.2.1	Superbatch Buffering	52
13.2.2	Accelerated Routines	52
13.2.3	Memory Management	52
13.3	Fortran–CUDA Interface	52
13.4	Multi-GPU Support	53
13.5	Verification	53
II Namelist Reference		
14	RUN_PARAMS	57
15	AMR_PARAMS	59

16	OUTPUT_PARAMS	61
17	INIT_PARAMS	63
18	REFINE_PARAMS	65
19	HYDRO_PARAMS	67
20	POISSON_PARAMS	69
21	PHYSICS_PARAMS	71
22	Example Namelist	73
22.1	Cosmological Simulation with Memory Balancing	73

III Build and Testing

23	Build Instructions	77
23.1	Prerequisites	77
23.2	Build	77
23.3	Compile-Time Options	77
24	Verification Tests	79
24.1	Test Configuration	79
24.2	Reference Values (nremap=5)	79
24.3	Running Tests	79
A	Modified Files Summary	81



Code Modifications

1

Quick Start Guide

This chapter provides a practical step-by-step guide to building cuRAMSES and running a cosmological zoom-in simulation. For detailed explanations of the underlying algorithms and parameters, refer to the subsequent chapters.

1.1 Prerequisites

- **Intel Fortran Compiler** (ifx or ifort) with MPI wrapper (mpiifx)
- **OpenMP** support (enabled via -qopenmp)
- **MUSIC** — Multi-Scale Initial Conditions generator (<https://bitbucket.org/ohahn/music>)
- **Python 3** with numpy (for utility scripts such as `create_pvar006.py`)

1.2 Building cuRAMSES

Build Commands

```
cd bin  
make clean  
make
```

The binary is produced as `bin/ramses_final3d`. Key compile-time options (-DNVAR=11, -DNPRE=8, -DLONGINT, etc.) are set in the Makefile. See Chapter 23 for details on compile-time flags.

1.3 Setting Up a Zoom-In Simulation

The following steps walk through a complete zoom-in cosmological simulation with star formation, targeting ~ 1 kpc resolution in a $10 h^{-1}$ Mpc box.

1.3.1 Step 1: Generate Initial Conditions with MUSIC

Create a MUSIC configuration file. The key settings for a zoom-in are `levelmin` (base grid), `levelmax` (finest zoom level), and `ref_extent` (zoom region fraction).

zoomin_10Mpc.conf (excerpt)

```
[setup]  
boxlength      = 10          # Box size in Mpc/h  
zstart         = 50          # Starting redshift  
levelmin       = 7           # Base grid: 128^3
```

```

levelmax          = 11           # Zoom finest: 2048^3 effective
ref_center        = 0.5, 0.5, 0.5
ref_extent        = 0.04, 0.04, 0.04
baryons           = yes
use_2LPT          = yes

[cosmology]
Omega_m           = 0.3111
Omega_L           = 0.6889
Omega_b           = 0.04
H0                = 67.66
sigma_8           = 0.8102
nspec              = 0.9665
transfer          = eisenstein

[output]
format             = grafic2
filename           = IC_zoomin

```

Run MUSIC:

Generate ICs

```
./MUSIC zoomin_10Mpc.conf
```

This produces IC_zoomin/level_007 through IC_zoomin/level_011, each containing GRAFIC2 binary files (ic_deltab, ic_velcx/y/z, ic_poscx/y/z, etc.).

1.3.2 Step 2: Create Zoom Geometry Scalar

The zoom geometry scalar (ic_pvar_00006) is a passive variable that marks which cells belong to the zoom region. This is essential when using ivar_refine=11 to control AMR refinement during initialization. See Chapter 10 for a detailed explanation.

Create pvar006

```
cd test_ksection
python3 create_pvar006.py
```

Edit the script's IC_DIR and LEVELS variables to match your IC directory.

1.3.3 Step 3: Create Run Directory

Setup Run Directory

```

mkdir -p run_zoomin
cd run_zoomin
ln -s ../IC_zoomin .
cp ../../bin/ramses_final3d .
cp ../../cosmo_zoomin_physics.nml namelist.nml

```

1.3.4 Step 4: Configure Namelist

The most important parameters to set correctly in the namelist:

Key namelist parameters

```

&RUN_PARAMS
ordering='ksection'          ! hierarchical domain decomposition
memory_balance=.true.         ! memory-weighted load balancing
nremap=5                      ! load balance every 5 coarse steps
/

&AMR_PARAMS
levelmin=7                    ! must match MUSIC levelmin
levelmax=20                   ! maximum AMR level allowed
ngridtot=200000000            ! total grids (see memory sizing)
nparttot=200000000            ! total particles
/

&REFINE_PARAMS
m_refine=13*8.                ! Lagrangian refinement threshold
ivar_refine=11                 ! use passive scalar for zoom mask
var_cut_refine=0.01             ! threshold for zoom scalar
mass_cut_refine=-1             ! set per IC level (see table below)
/

&INIT_PARAMS
filetype='grafic'
initfile(1)='IC_zoomin/level_007'
initfile(2)='IC_zoomin/level_008'
initfile(3)='IC_zoomin/level_009'
initfile(4)='IC_zoomin/level_010'
initfile(5)='IC_zoomin/level_011'
/

```

1.3.5 Step 5: Run

Launch Simulation

```
mpirun -np 12 ./ramses_final3d namelist.nml 2>&1 | tee run.log
```

1.4 Key Parameters Quick Reference

Parameter	Recommended	Purpose
ordering	'ksection'	Hierarchical domain decomposition
memory_balance	.true.	Memory-weighted load balancing
nremap	5	Load balance frequency (optimal)
ivar_refine	11 (=NVAR)	Passive scalar index for zoom mask
var_cut_refine	0.01	Threshold for zoom scalar refinement

Parameter	Recommended	Purpose
mass_cut_refine	see table	Particle mass filter for cpu_map2
ngridtot	see sizing	Total grids; keep virtual mem < CommitLimit
nparttot	see sizing	Total particles across all ranks
levelmin	IC levelmin	Must match MUSIC base level
levelmax	14–20	Maximum AMR refinement

1.4.1 mass_cut_refine Reference

The `mass_cut_refine` parameter filters out heavy (background) dark matter particles from the `cpu_map2` density computation in `rho_fine`. Set it to a value between the zoom-region particle mass and the coarser-level particle mass. Recommended values by IC finest level:

IC finest level	Zoom m_{DM}	Coarse m_{DM}	<code>mass_cut_refine</code>
9	$\sim 4.6 \times 10^7$	$\sim 3.7 \times 10^8$	1.0×10^8
10	$\sim 5.8 \times 10^6$	$\sim 4.6 \times 10^7$	1.5×10^7
11	$\sim 7.2 \times 10^5$	$\sim 5.8 \times 10^6$	2.0×10^6
12	$\sim 9.0 \times 10^4$	$\sim 7.2 \times 10^5$	2.5×10^5
13	$\sim 1.1 \times 10^4$	$\sim 9.0 \times 10^4$	3.0×10^4

Units

Masses are in M_{\odot}/h for a $10 h^{-1}$ Mpc box with Planck 2018 cosmology ($\Omega_m = 0.3111$, $\Omega_b = 0.04$, $h = 0.6766$). The exact values depend on the box size and cosmological parameters.

1.5 Troubleshooting

Error / Symptom	Solution
No more free memory	Increase <code>ngridtot</code> (or <code>ngridmax</code>) in <code>&AMR_PARAMS</code> .
No more free memory for particles	Increase <code>nparttot</code> (or <code>npartmax</code>) in <code>&AMR_PARAMS</code> .
Process killed (SIGNAL 9, OOM)	<code>ngridtot</code> is too large for available RAM. RAMSES allocates the full array at startup (virtual memory). Check <code>CommitLimit</code> in <code>/proc/meminfo</code> and reduce <code>ngridtot</code> so that total virtual memory stays below it.
Background AMR explosion (entire box refined, memory exhausted)	Verify <code>ivar_refine=11</code> and that <code>ic_pvar_00006</code> files exist in each IC level directory. Without the zoom geometry scalar, <code>cpu_map2</code> marks the entire domain for refinement.

Error / Symptom	Solution
Refinement does not extend beyond IC levels	Check that <code>m_refine</code> has enough entries (one per extra level above IC finest) and that <code>mass_cut_refine</code> is set correctly to exclude heavy background particles.
Very slow Poisson solver	Increase <code>cg_levelmin</code> in <code>&POISSON_PARAMS</code> to switch from multigrid to conjugate-gradient at the finest levels. Typical: <code>cg_levelmin=14</code> .

2

K-Section Ordering

2.1 Overview

The **k-section ordering** replaces the standard Hilbert curve domain decomposition with a hierarchical spatial bisection tree. Unlike the Hilbert curve approach that relies on a 1D space-filling curve, the k-section ordering partitions the 3D domain using recursive bisection along each coordinate axis, producing a balanced k-ary tree of spatial subdomains.

Key Advantage

The k-section tree enables **hierarchical MPI exchange** where communication follows the tree structure, achieving $O(\sum k_l)$ messages per exchange instead of $O(N_{\text{cpu}})$ all-to-all communication.

Figure 2.2 illustrates the progressive nature of the k-section tree construction. Starting from the full simulation domain, the tree is built by recursively bisecting (or k -secting) along successive coordinate axes. In this example with $k_1 \times k_2 \times k_3 = 3 \times 2 \times 2 = 12$ CPUs, the decomposition proceeds in three stages:

1. **Level 1** ($k_1 = 3$): The domain is split into 3 slabs along the x -axis.
2. **Level 2** ($k_2 = 2$): Each slab is bisected along the y -axis, producing 6 subdomains.
3. **Level 3** ($k_3 = 2$): Each subdomain is bisected along the z -axis, yielding the final 12 subdomains.

The split positions at each level are determined by the load-balancing histogram to ensure equal computational cost per subdomain. The color hue in Figure 2.2 encodes the x -slab membership (red/green/blue), with saturation and brightness variations indicating the y and z subdivisions, respectively.

2.2 Hierarchical MPI Exchange

Two exchange routines are provided in `patch/cuda/ksection.f90`:

- `ksection_exchange_dp` — exclusive exchange (each item → exactly 1 destination CPU)
- `ksection_exchange_dp_overlap` — overlap exchange (items routed by spatial bounding box)

2.2.1 Exclusive Exchange

Each data item has a known destination CPU. The algorithm performs level-by-level correspondent exchange through the k-section tree, using MPI tags 100–300+level.

Domain Decomposition $3 \times 3 \times 2 = 18$ subboxes
 4096^3 TSC density | $\log_{10}(1+\Sigma)$ projected column density

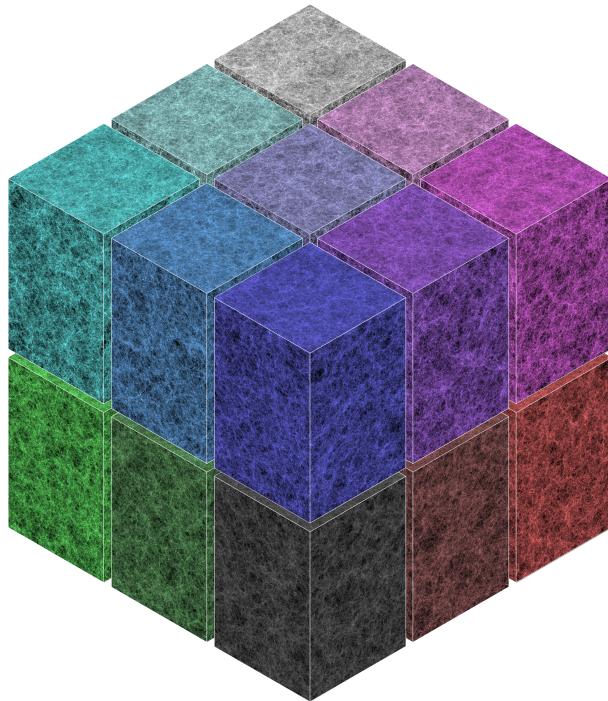


Figure 2.1: Isometric view of k-section domain decomposition with a $3 \times 3 \times 2 = 18$ uniform partition. Each subbox surface shows the projected column density $\log_{10}(1 + \Sigma)$ from a 4096^3 TSC density field. Colors encode spatial position: R → x , G → y , B → z , providing a smooth gradient across adjacent subboxes.

2.2.2 Overlap Exchange

Items have spatial extent (bounding box) and may need to reach multiple CPUs. The tree walk determines all destination CPUs whose spatial domain overlaps the item's bounding box. MPI tags 400–500+level are used.

2.2.3 Periodic Boundary Conditions

The optional `periodic=.true.` parameter enables handling of items that wrap around the domain boundary $[0, \text{scale}]$. For each wrapping dimension, $2^{n_{\text{wrap}}} - 1$ shifted copies are generated via bitmask subset enumeration before the tree walk.

2.3 Tree Navigation

The arrays `ksec_cpumin/cpumax` and `ksec_cpu_path` (set in `build_ksection` and at restart via `rebuild_ksec_cpuranges`) enable each CPU to navigate the tree efficiently.

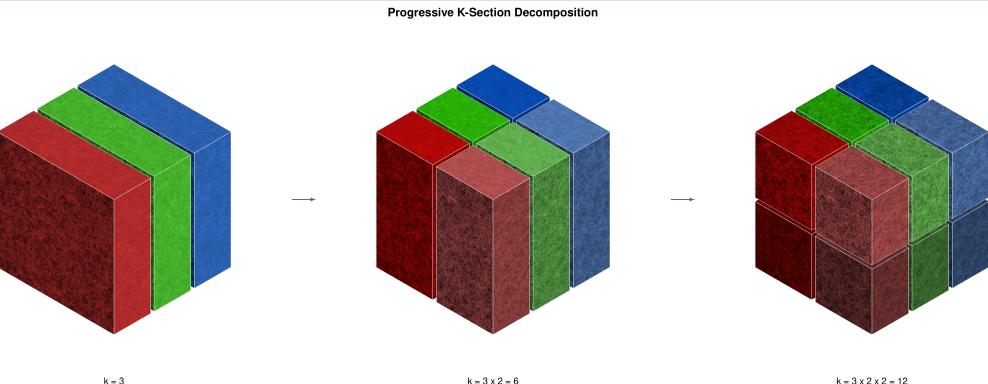


Figure 2.2: Progressive k-section decomposition for $k_1 \times k_2 \times k_3 = 3 \times 2 \times 2 = 12$ CPUs. **Left:** first split into 3 slabs along x . **Center:** each slab bisected along y ($3 \times 2 = 6$ subdomains). **Right:** final bisection along z ($3 \times 2 \times 2 = 12$ subdomains). Each face shows projected column density $\log_{10}(1 + \Sigma)$ from a 4096^3 TSC density field.

2.4 Verification

The exchange routines are tested in `test_ksection_exchange` with 4 test cases:

1. Exclusive point exchange
2. Overlap point exchange
3. Full overlap exchange
4. Periodic overlap exchange (20% radius items)

All tests pass.

3

Ghost Zone Exchange via K-Section

3.1 AMR Ghost Zones

The standard RAMSES ghost zone exchange uses MPI_ISEND/IRecv with all-to-all communication patterns. This was replaced with k-section tree-routed exchange for the `ordering='ksection'` mode.

3.1.1 Modified File

`patch/cuda/virtual_boundaries.kjhan.f90`

3.1.2 Subroutines

Four k-section variants were implemented:

Subroutine	Direction	Data Type
<code>make_virtual_fine_dp_ksec</code>	Forward	<code>real(dp)</code>
<code>make_virtual_fine_int_ksec</code>	Forward	<code>integer</code>
<code>make_virtual_reverse_dp_ksec</code>	Reverse (+= accumulate)	<code>real(dp)</code>
<code>make_virtual_reverse_int_ksec</code>	Reverse (+= accumulate)	<code>integer</code>

3.1.3 Data Packing

Each emission grid is packed as:

$$\text{sendbuf}(1:\text{twotondim} + 2, i) = [\underbrace{c_1, c_2, \dots, c_8}_{\text{cell data}}, \underbrace{\text{myid}}_{\text{sender}}, \underbrace{i}_{\text{index}}]$$

The metadata (sender ID, emission index) allows the receiver to scatter data to the correct reception grid without requiring a priori knowledge of the communication pattern.

3.1.4 Dispatch

Dispatch is automatic via:

```
if (ordering=='ksection') then
    call make_virtual_fine_dp_ksec(xx, illevel)
    return
end if
```

3.1.5 Bulk Exchange

Four bulk variants exchange all columns of a 2D array (e.g., `uold`, `f`, `unew`) in a single `ksection_exchange_dp` call:

Subroutine	Description
<code>make_virtual_fine_dp_bulk</code>	Forward bulk exchange
<code>make_virtual_fine_dp_bulk_ksec</code>	Forward bulk (ksection impl.)
<code>make_virtual_reverse_dp_bulk</code>	Reverse bulk exchange
<code>make_virtual_reverse_dp_bulk_ksec</code>	Reverse bulk (ksection impl.)

Buffer layout for `ncols` variables:

$$\text{sendbuf}((v-1) \cdot 2^3 + j, \text{idx}) = \text{xx}(\text{icell}, v) \quad v = 1 \dots \text{ncols}, j = 1 \dots 8$$

plus 2 metadata words (sender ID, index). This reduces MPI exchanges from N_{var} per level to 1 per array.

3.2 build_comm via K-Section

The `build_comm` subroutine's `MPI_ALLTOALL` + `MPI_ISEND`/`IRecv` pattern was also replaced with `ksection_exchange_dp`.

3.3 MPI_ALLTOALL Replacements

Additional `MPI_ALLTOALL` calls in the following files were replaced with k-section exchange:

- `particle_tree.kjhan.f90`
- `init_part.f90`
- `multigrid_fine_commons.f90` (in `build_parent_comms_mg`)

3.4 Pre-Allocated Buffer Pool

Per-level small arrays (child count, peer list, MPI requests) were converted to save variables to eliminate ~ 100 allocations/deallocations per call. The `peer_recv` buffer uses grow-only capacity, and the first dimension must exactly match `nprops` for MPI stride correctness.

4

Multigrid Poisson K-Section Communication

4.1 Motivation

The multigrid Poisson solver (`poisson-mg`) accounts for 29–41% of total runtime. The original MPI communication used `MPI_ISEND`/`IRECV` with all-to-all patterns across all CPUs.

4.2 Implementation

Four k-section variants were added to `patch/cuda/multigrid_fine_commons.f90`:

Subroutine	Direction	Data Type
<code>make_virtual_mg_dp_ksec</code>	Forward	<code>real(dp)</code>
<code>make_virtual_mg_int_ksec</code>	Forward	<code>integer</code>
<code>make_reverse_mg_dp_ksec</code>	Reverse (+= accumulate)	<code>real(dp)</code>
<code>make_reverse_mg_int_ksec</code>	Reverse (+= accumulate)	<code>integer</code>

4.2.1 Key Differences from AMR Exchange

The MG communication uses different data structures from the standard AMR exchange:

Aspect	AMR	Multigrid
Active grids	<code>active(ilevel)</code>	<code>active_mg(myid,ilevel)</code>
Reception	<code>reception(icpu,ilevel)</code>	<code>active_mg(icpu,ilevel)</code>
Emission	<code>emission(icpu,ilevel)</code>	<code>emission_mg(icpu,ilevel)</code>
Data arrays	<code>xx(igrid+iskip)</code>	<code>active_mg%u(icell,ivar)</code>
Indexing	Global grid index	Local offset in <code>active_mg</code>

4.2.2 Forward Exchange

1. Pack: for each `emission_mg(icpu)%igrid(i)`, gather `two_tondim` values from `active_mg(myid)%u` + metadata (`sender_id, index`)
2. Exchange via `ksection_exchange_dp`
3. Scatter: use metadata to write into `active_mg(sender)%u(ridx + step, ivar)`

4.2.3 Reverse Exchange (Accumulation)

1. Pack: for each remote `active_mg(icpu)%u(i + step, ivar) + metadata`
2. Exchange via `ksection_exchange_dp`
3. Accumulate: `active_mg(myid)%u(emission_mg(sender)%igrid(ridx) + step, ivar) += recvbuf`

5

Morton Key Octree

5.1 Overview

The `nbor` array (6 neighbor pointers per grid) was replaced with a Morton key hash table for $O(1)$ neighbor lookup. This saves $6 \times 8 \times N_{\text{gridmax}}$ bytes of memory.

5.1.1 Modified Files

- `patch/oct_tree/morton_keys.f90` — Morton key computation
- `patch/oct_tree/morton_hash.f90` — Hash table and helper functions
- `patch/oct_tree/morton_init.f90` — Initialization and verification
- `patch/oct_tree/refine_utils.f90` — Hash table maintenance
- `patch/oct_tree/nbors_utils.kjhan.f90` — Neighbor lookup functions

5.2 Morton Key

A 64-bit Morton key is computed by interleaving the 3D integer coordinates (21 bits per dimension):

$$\text{key} = \text{interleave}(\lfloor x_g \cdot 2^{l-1} \rfloor, \lfloor y_g \cdot 2^{l-1} \rfloor, \lfloor z_g \cdot 2^{l-1} \rfloor)$$

5.3 Hash Table

A per-level open-addressing hash table with linear probing and power-of-2 capacity maps Morton keys to grid indices. The table is maintained in `make_grid_coarse/fine` and `kill_grid`, with a full rebuild after each time step.

5.4 Neighbor Lookup

Two helper functions in the `morton_hash` module:

- `morton_nbor_grid(igrid, ilevel, j)` — returns `son(nbor(igrid,j))` equivalent
 - `morton_nbor_cell(igrid, ilevel, j)` — returns `nbor(igrid,j)` equivalent
- Direction convention: $j = 1:-x, 2:+x, 3:-y, 4:+y, 5:-z, 6:+z$.

5.5 nbor Array Removal (Phase 4)

The `nbor` array is allocated as `allocate(nbor(1:1,1:1))` (minimum size to avoid compilation errors). All code paths use Morton lookup exclusively.

6

Memory-Based Load Balancing

6.1 Motivation

Standard RAMSES load balancing distributes cells evenly across CPUs, but cells with many particles consume significantly more memory. Memory-based balancing weights each cell by its memory footprint.

6.2 Cost Function

$$\text{cell_cost} = \frac{\text{mem_weight_grid}}{\text{twotondim}} + \text{numbp}(\text{igrid}) \times \frac{\text{mem_weight_part}}{\text{twotondim}}$$

where:

- `mem_weight_grid` = 270 (default) — memory per grid in dp-equivalents
- `mem_weight_part` = 12 (default) — memory per particle in dp-equivalents

6.3 Implementation Details

- All histogram variables use 64-bit integers (`integer(i8b)`) with `MPI_INTEGER8`
- `numbp` is synchronized for virtual/reception grids before cost computation, then restored afterwards
- The `numbp` restore uses a save/restore pattern to avoid breaking the particle tree

6.4 Parameters

Controlled by three namelist parameters in `&RUN_PARAMS`:

- `memory_balance = .true.` — enable memory-based balancing
- `mem_weight_grid = 270` — grid memory weight
- `mem_weight_part = 12` — particle memory weight

7

Memory Savings: Large Array Optimization

7.1 Overview

Several large arrays were eliminated or converted to on-demand allocation to reduce steady-state memory usage by \sim 960 MB (for `ngridmax=5M`).

Array	Strategy	Savings
<code>hilbert_key</code>	<code>allocate(1:1)</code> for ksection	\sim 640 MB
<code>bisec_ind_cell + cell_level</code>	On-demand alloc/dealloc	\sim 320 MB
<code>defrag_map</code>	Local scratch during defrag	minor
<code>nbor</code>	<code>allocate(1:1, 1:1)</code> (Morton)	\sim 240 MB

8

IC Reading with Stream Access

8.1 Motivation

Sequential Fortran I/O requires reading all preceding planes to reach a target plane, which is $O(n^2)$ for large files. Stream access enables direct byte-offset seeks.

8.2 Implementation

Fortran 2003 ACCESS='STREAM' is used with computed byte offsets:

```
hdr_bytes = 52 + (i3-1)*plane_bytes + 5
plane_bytes = n1*n2*4 + 8 ! data + 2 record markers
```

Applied to hydro IC (deltab, velocity, temperature), particle velocity and position files. Only for multiple=.false. mode.

8.2.1 Modified Files

- patch/Horizon5-master-2/init_flow_fine.f90
- init_part.f90

9

Load Balance Profiling and Tuning

9.1 Internal Timing

Detailed timing breakdown was added to `load_balance` in `patch/cuda/load_balance.kjhan.f90`:

Section	Description	Typical (s/step)
<code>numbp_sync</code>	MPI sync of <code>numbp</code> for virtual grids	0.8–1.0
<code>cmp_new_cpu_map</code>	Build ksection + compute new map	0.4–0.6
<code>expand_pass</code>	<code>build_comm</code> + <code>make_virtual</code> loop	0.8–1.5
<code>grid_migration</code>	Linked-list reconnection	< 0.01
<code>allreduce+cpumap_update</code>	<code>MPI_ALLREDUCE</code> × 4 + <code>cpu_map</code>	2.3–3.3
<code>shrink_pass</code>	<code>flag_fine</code> + <code>build_comm</code> loop	0.4–1.0

9.2 nremap Tuning

The `nremap` parameter controls load balancing frequency (every N coarse steps). Testing with 200M particles on 12 ranks showed:

nremap	Total (s)	Loadbal (s)	Speedup	Note
1	303.8	64.4 (21.2%)	—	Baseline
3	269.9	24.7 (9.1%)	1.13×	
5	249.8	15.7 (6.3%)	1.22×	Optimal
10	258.6	11.6 (4.5%)	1.17×	Imbalance grows

Default Setting

`nremap=5` is set as the default. All four configurations produce **bit-identical** results: `econs=3.77E-03, epot=-1.88E-06, ekin=1.23E-06` at step 10.

9.3 Min/Max Memory Reporting

The `writemem_minmax` subroutine prints per-step min/max memory usage across all MPI ranks.

10

Zoom-In Simulation Setup

10.1 Overview

This chapter describes how to configure and run a cosmological zoom-in simulation with cuRAMSES, targeting ~ 1 kpc physical resolution within a selected sub-region of a larger cosmological volume. The setup uses:

- **Box size:** $10 h^{-1}$ Mpc (comoving)
- **Zoom region:** $\sim 1.25 h^{-1}$ Mpc (centered, with padding from MUSIC)
- **IC levels:** 7–11 (generated by MUSIC), corresponding to base 128^3 up to effective 2048^3
- **AMR levels:** up to 14 (adaptive refinement beyond IC levels)
- **Physics:** hydrodynamics, gravity, radiative cooling, star formation, and AGN feedback

The main challenge in zoom-in simulations is preventing the AMR machinery from refining the entire box (which would exhaust memory). This is solved by using a passive scalar variable as a *zoom geometry mask* that restricts refinement to the zoom region during initialization.

10.2 Initial Conditions

10.2.1 MUSIC Configuration

Initial conditions are generated by the MUSIC code with a multi-level zoom setup. The MUSIC configuration file specifies the base resolution, zoom region, and cosmological parameters.

`zoomin_10Mpc.conf`

```
[setup]
boxlength      = 10          # Box size in Mpc/h
zstart         = 50          # Starting redshift
levelmin       = 7           # Base grid: 2^7 = 128^3
levelmin_TF    = 7           # Transfer function grid
levelmax       = 11          # Zoom finest: 2^11 = 2048^3
                           # effective
padding        = 8
overlap        = 4
ref_center     = 0.5, 0.5, 0.5
ref_extent     = 0.04, 0.04, 0.04
align_top      = yes
baryons        = yes
use_2LPT       = yes
periodic_TF   = yes

[cosmology]
# Planck 2018 (TT, TE, EE+lowE+lensing)
```

```

Omega_m          = 0.3111
Omega_L          = 0.6889
Omega_b          = 0.04
H0               = 67.66
sigma_8          = 0.8102
nspec            = 0.9665
transfer         = eisenstein

[output]
format           = grafic2
filename         = IC_zoomin

```

10.2.2 Generated IC Structure

MUSIC produces a directory hierarchy with one sub-directory per refinement level:

Level	Grid Size	Resolution	Coverage
7	128^3	$\sim 78 \text{ kpc}/h$	Entire box
8	variable	$\sim 39 \text{ kpc}/h$	Zoom region + padding
9	variable	$\sim 20 \text{ kpc}/h$	Zoom region + padding
10	variable	$\sim 10 \text{ kpc}/h$	Zoom region + padding
11	variable	$\sim 5 \text{ kpc}/h$	Zoom region (finest IC)

Each level directory contains GRAFIC2 binary files: `ic_deltab` (baryon overdensity), `ic_velcx/y/z` (baryon velocities), `ic_velcdx/y/z` (dark matter velocities), `ic_poscdx/y/z` (dark matter position offsets), and `ic_refmap` (refinement map, level 7 only).

10.3 Zoom Geometry Scalar (`ic_pvar_00006`)

Important

This is the single most critical aspect of the zoom-in setup. Incorrect configuration of the zoom geometry scalar will cause the entire box to be refined, leading to immediate memory exhaustion.

10.3.1 The Problem

In a standard RAMSES simulation, the `init_refmap` subroutine (called when `ivar_refine=0`) reads the `ic_refmap` file and sets `cpu_map2=1` for cells that should be refined. However, the default `flag_utils` logic using `cpu_map2` applies refinement uniformly wherever the map is set, without distinguishing between zoom and background regions at higher AMR levels. This causes the entire box to be recursively refined, rapidly exhausting memory.

10.3.2 The Solution: Passive Scalar Mask

The solution uses the 6th passive scalar variable (`ic_pvar_00006`) as a zoom geometry indicator. With `NVAR=11` (5 hydro + 1 metal + 5 passive), the 6th passive scalar corresponds to hydro variable index 11. Setting `ivar_refine=11` activates a different refinement criterion in `flag_utils`:

```

! In flag_utils.kjhan.f90 (during init):
if(ivar_refine > 0) then
  do i=1,ncell
    ok(i) = ok(i) .or. &
      (uold(ind_cell(i),ivar_refine) / uold(ind_cell(i)
        ,1) &
       > var_cut_refine)
  end do
end if

```

This checks whether the passive scalar (divided by density for the conserved-to-primitive conversion) exceeds `var_cut_refine` (typically 0.01). Cells in the zoom region have a scalar value of 1.0 and pass the test; background cells have 0.0 and are skipped.

10.3.3 How It Works at Each Stage

1. **During initialization** (`init=.true.`): When `ivar_refine` is nonzero, `init_refmap` is *not* called (line 31 of `amr/init_refine.f90`):

```

if(ivar_refine==0) call init_refmap

```

Instead, `init_flow` loads the hydro IC (including `ic_pvar_00006`) and the refinement flag is set by the passive scalar criterion: $uold(cell,11)/uold(cell,1) > 0.01$.

2. **After initialization** (`init=.false.`): The refinement criterion switches to the standard density-based Lagrangian criterion (`m_refine`). The `cpu_map2` array is set by `rho_fine` with the `mass_cut_refine` parameter filtering out heavy background dark matter particles, ensuring that only the zoom region (populated by light, high-resolution particles) triggers further refinement.

10.3.4 Creating `ic_pvar_00006`

The Python script `test_ksection/create_pvar006.py` generates the zoom geometry scalar for each IC level:

- **Level 7** (base): Reads `ic_refmap` from the level 7 directory. Nonzero entries (zoom cells) are set to 1.0; zero entries (background) are set to 0.0. The result is written in GRAFIC2 format as `ic_pvar_00006`.
- **Levels 8+** (zoom sub-levels): All cells are set to 1.0, since the entire grid at these levels lies within the zoom region by construction.

`create_pvar006.py` (core logic)

```

for level in LEVELS:
  if level == base_level:
    # Read ic_refmap, convert: nonzero -> 1.0, zero -> 0.0
    refmap = read_grafic2_data(refmap_file, n1, n2, n3)
    pvar = np.where(refmap != 0, 1.0, 0.0).astype(np.float32)
  else:
    # Zoom sub-levels: all cells = 1.0
    pvar = np.ones((n3, n2, n1), dtype=np.float32)
    write_grafic2(pvar_file, header_bytes, pvar)

```

Editing the Script

Before running, edit the `IC_DIR` variable in `create_pvar006.py` to point to your IC directory, and adjust `LEVELS` to match the levels generated by MUSIC.

10.3.5 mass_cut_refine

After initialization, `rho_fine` computes the density field used to set `cpu_map2` for refinement flagging. The `mass_cut_refine` parameter acts as a mass filter:

```
! In rho_fine.f90:
if(mass_cut_refine > 0.0) then
    do j = 1, np
        if(ttt(j) == 0d0) then           ! dark matter only
            ok(j) = ok(j) .and. mmm(j) < mass_cut_refine
        endif
    end do
endif
```

Particles heavier than `mass_cut_refine` are excluded from the density computation, so only high-resolution zoom particles contribute to the refinement map. Set `mass_cut_refine` to a value between the zoom-region DM particle mass and the next-coarser level's DM particle mass. Refer to the table in Chapter 1 for recommended values.

10.4 Memory Considerations

10.4.1 ngridtot Sizing

RAMSES allocates all grid-related arrays at startup based on `ngridmax` (per-CPU maximum grids). When `ngridtot` is specified, `ngridmax` is computed as `ngridtot/ncpu`. The critical constraint is:

Important

The **total virtual memory** allocated by all MPI ranks on a node must not exceed the kernel's `CommitLimit`. Exceeding this causes the OOM killer to terminate the process (SIGNAL 9).

10.4.2 CommitLimit

The Linux kernel's `CommitLimit` determines the maximum total virtual memory that can be allocated:

$$\text{CommitLimit} = \text{RAM} \times \frac{\text{overcommit_ratio}}{100} + \text{swap}$$

Check the current value:

Check CommitLimit

```
grep CommitLimit /proc/meminfo
```

On systems without swap (common for HPC nodes), `CommitLimit = RAM × overcommit_ratio/100`. The default `overcommit_ratio` is 50, so a 256 GB node has ~128 GB `CommitLimit`.

10.4.3 Virtual vs. Physical Memory

RAMSES allocates the full `ngridmax`-sized arrays at startup, consuming virtual memory immediately. Physical (resident) memory grows as grids are actually created during the simulation. The key arrays that dominate virtual memory usage are:

Array	Size	Per-grid bytes
<code>uold(1:ncell,1:nvar)</code>	$8 \times 8 \times nvar$	$8 \times 8 \times 11 = 704$
<code>unew(1:ncell,1:nvar)</code>	same	704
<code>f(1:ncell,1:3)</code>	$8 \times 8 \times 3$	192
<code>son/father/next/prev</code>	8×8 each	$4 \times 64 = 256$
<code>phi/rho</code>	8×8 each	$2 \times 64 = 128$
<code>flag1/flag2/cpu_map/map2</code>	4×8 each	$4 \times 32 = 128$

As a rough estimate, each grid (oct = 8 cells) requires ~2–3 kB of virtual memory across all arrays. For `ngridmax`=25 million (i.e., `ngridtot`=200M with 8 CPUs), each rank allocates ~50–75 GB of virtual memory. Ensure that `ncpu × per_rank_virtual < CommitLimit`.



11.1 Overview

Standard RAMSES writes one binary file per MPI process per data type (AMR, hydro, gravity, particles), producing $O(N_{\text{cpu}} \times 4)$ files per snapshot. This creates file management challenges at high core counts and requires the same number of MPI processes for restart.

The HDF5 I/O module replaces the per-CPU binary output with a single HDF5 file per snapshot, using MPI parallel I/O (collective hyperslab writes) for performance. Key benefits:

- **Single file:** one `data_NNNNN.h5` per snapshot instead of thousands of files
- **Parallel I/O:** MPI-IO backend ensures scalable write/read performance
- **Self-describing:** HDF5 groups and attributes contain all metadata
- **Cross-format:** `informat` and `outformat` can differ, enabling binary-to-HDF5 conversion

Compilation

HDF5 support requires compilation with `make HDF5=1`. The HDF5 library must be built with `-enable-parallel` and `-enable-fortran` for the Intel ifx compiler. See Section 11.2 for build instructions.

11.2 Building HDF5

The HDF5 library must be compiled from source with parallel (MPI-IO) and Fortran support matching the Intel compiler used for RAMSES.

Build HDF5 from source

```
wget https://github.com/HDFGroup/hdf5/releases/download/\
      hdf5-1.14.5/hdf5-1.14.5.tar.gz
tar xf hdf5-1.14.5.tar.gz && cd hdf5-1.14.5
CC=mpiicc FC=mpiifx ./configure \
  --enable-fortran --enable-parallel \
  --prefix=$HOME/local/hdf5
make -j8 && make install
```

The key configure flags:

- `--enable-parallel`: MPI-IO support for collective parallel reads/writes
 - `--enable-fortran`: generates `hdf5.mod` (Fortran module) compatible with `ifx`
- Then build RAMSES with the `HDF5=1` flag:

Build RAMSES with HDF5

```
cd bin
make clean
make HDF5=1
```

The Makefile conditionally adds:

```
HDF5_DIR = $(HOME)/local/hdf5
FFLAGS += -DHDF5 -I$(HDF5_DIR)/include
LIBS += -L$(HDF5_DIR)/lib -lhdf5_fortran -lhdf5 -lz \
        -Wl,-rpath,$(HDF5_DIR)/lib
```

Without `HDF5=1`, the code compiles identically to before (all HDF5 code is guarded by `#ifdef HDF5`).

11.3 Namelist Parameters

Two parameters in `&OUTPUT_PARAMS` control the I/O format:

Parameter	Type	Default	Description
<code>outformat</code>	char(10)	'original'	Output format: 'original' or 'hdf5'
<code>informat</code>	char(10)	'original'	Restart format: 'original' or 'hdf5'

Example namelist

```
&OUTPUT_PARAMS
noutput=1
aout=1.0
foutput=5
outformat='hdf5'      ! write HDF5 snapshots
/
```

For restart from an HDF5 checkpoint:

Restart from HDF5

```
&RUN_PARAMS
nrestart=1
/
&OUTPUT_PARAMS
informat='hdf5'      ! read HDF5 checkpoint
outformat='hdf5'      ! continue writing HDF5
/
```

11.4 HDF5 File Structure

Each snapshot produces a single file `output_NNNNNN/data_NNNNNN.h5` with the following group hierarchy:

```

data_NNNNN.h5
  /header/          (attrs: ncpu, ndim, nlevelmax, nstep,
                    boxlen, time, aexp, H0, omega_m, ...)
    tout[noutput]   (dataset)
    aout[noutput]   (dataset)
    dtold[nlevelmax] (dataset)
    dtnew[nlevelmax] (dataset)
  /domain/          (ksection tree or bound_key)
  /coarse/
    son[ncoarse]    (dataset: integer)
    cpu_map[ncoarse] (dataset: integer)
  /amr/level_LL/
    xg[ngrid_total x 3]      (grid centres)
    son[ngrid_total x twotondim] (son indices)
    cpu_map[ngrid_total x twotondim]
    nbor[ngrid_total x twondim] (Morton-computed)
    ngrid_per_cpu[ncpu]
  /hydro/level_LL/
    uold[ngrid_total x twotondim x nvar]
  /gravity/level_LL/
    phi[ngrid_total x twotondim]
    f[ngrid_total x twotondim x ndim]
  /particles/
    x[npart_total x ndim]     (positions)
    v[npart_total x ndim]     (velocities)
    m[npart_total]           (masses)
    id[npart_total]          (particle IDs)
    level[npart_total]       (AMR level)
    tp[npart_total]          (birth time)
    zp[npart_total]          (metallicity)
    npart_per_cpu[ncpu]
  /sinks/
    idsink, msink, xsink, vsink, tsink, ...

```

11.4.1 Parallel Write Strategy

Each AMR level is written using collective MPI-IO hyperslabs:

1. `MPI_Allgather(ngrid_local)` to compute per-CPU offsets
2. Each CPU writes its portion via an HDF5 hyperslab selection
3. Collective I/O ensures efficient parallel access to the shared file

Particles use the same hyperslab pattern: `MPI_Allgather(npart_local)` for offsets, then collective write.

11.5 Implementation Files

File	Description
<code>patch/cuda/ramses_hdf5_io.f90</code>	HDF5 wrapper module (create, open, write, read helpers)
<code>patch/cuda/backup_hdf5.f90</code>	HDF5 output: <code>dump_all_hdf5()</code>
<code>patch/cuda/restore_hdf5.f90</code>	HDF5 restart: <code>restore_amr_hdf5()</code> , etc.

11.5.1 Dispatch

The HDF5 output is dispatched from `dump_all` in `output_amr.kjhan.f90`:

```
#ifdef HDF5
if(outformat == 'hdf5') then
    call dump_all_hdf5(filedir, nchar)
    goto 998 ! skip binary output
end if
#endif
```

For restart, each `init_*` subroutine checks `informat`:

```
#ifdef HDF5
if(informat == 'hdf5') then
    call restore_amr_hdf5()
    return
end if
#endif
```

11.6 Verification

The HDF5 output was tested with a 5-step cosmological simulation (12 MPI ranks, `levelmin=8`, `levelmax=10`, 200M particles):

Metric	Value
Output file	<code>data_00001.h5</code> (3.65 GB)
Step count	5
econs at step 5	4.85E-03
I/O time	4.1 s (16.1% of total)
Total runtime	20.6 s

11.7 Variable-NCPU Restart: Distributed Grid Creation

When restarting an HDF5 checkpoint with a different number of MPI ranks (`ncpu_file` \neq `ncpu`), the original implementation allocated *all* grids from the file on *every* rank. For a file with 11.17 M grids, this required `ngridmax` \geq 11.17 M per rank, consuming \sim 20 GB per rank. With 32 ranks the total memory exceeded 640 GB, causing out-of-memory failures on systems with \leq 560 GB.

11.7.1 Distributed Algorithm

The new implementation creates only the grids each rank actually needs. For each level $L = 1$ to `nlevelmax_file`:

1. **Phase 1 — Active grid creation.** All ranks read the file data for level L (grid positions and `son_flag`), but each rank only allocates grids whose father cell satisfies `cpu_map(father_cell) == myid`. For $L \geq 2$, the father grid is located via Morton hash lookup at level $L - 1$; if the lookup returns zero (father not present on this rank), the grid is not local and is skipped. Active grids have their `xg`, `flag1`, `cpu_map`, `cpu_map2`, `father`, and `son` set from the file data, and are inserted into the Morton hash table.

2. **Phase 2 — Virtual grid creation.** Virtual (ghost) grids are created using RAMSES's existing refinement infrastructure:

- $L = 1$: `flag_coarse` → `refine_coarse`
- $L \geq 2$: `refine_fine`($L - 1$)

Since active grids already set `son(father_cell) > 0`, `refine_fine` skips those cells and only creates virtual grids where `flag1 == 1` and `son == 0`. The flag `balance = .true.` causes `make_grid_fine` to skip hydro variable interpolation (line 791 of `refine_utils.f90`).

3. **Communication setup.** `build_comm`(L) establishes emission/reception arrays, followed by `make_virtual_fine_int` exchanges for `flag1`, `cpu_map`, and `cpu_map2`. This ensures the next level's `refine_fine` has correct flag values on virtual grids.

After all levels, hydro and Poisson data are restored from the file (same read-all-scatter pattern) and exchanged to virtual grids via `make_virtual_fine_dp`.

11.7.2 Key Design Decisions

- **Father lookup as filter.** For $L \geq 2$, the Morton hash lookup at level $L - 1$ serves as an implicit ownership test: if the father grid does not exist locally, the child grid cannot be active on this rank.
- **Hash table sizing.** The per-level hash table is sized for local grids: `max(4 * (ngrid_total / ncpu + 1), 16)` instead of `2 * ngrid_total`, reducing hash table memory proportionally.
- **No changes to other files.** `refine_coarse`, `refine_fine`, `build_comm`, `authorize_fine`, and `make_grid_fine` all work unmodified. The same-ncpu restart path is also unchanged.

11.7.3 Verification

A 12-CPU HDF5 checkpoint (`levelmin=8`, `levelmax=10`, 200M particles) was restarted with 8 CPUs. The reference is the same checkpoint restarted with 12 CPUs (same-ncpu path).

Metric	Reference (12→12)	Distributed (12→8)
ngrid per rank	2,396,745	352,379
Memory per rank	8.0 GB	6.9 GB
Grid reduction	—	6.8×

Step	econs		epot		ekin	
	Ref	Dist	Ref	Dist	Ref	Dist
6	8.18E-03	8.17E-03	-1.01E-06	-1.01E-06	6.49E-07	6.49E-07
7	7.26E-03	7.25E-03	-1.26E-06	-1.26E-06	8.17E-07	8.17E-07
8	6.43E-03	6.42E-03	-1.48E-06	-1.48E-06	9.63E-07	9.63E-07
9	5.79E-03	5.79E-03	-1.68E-06	-1.68E-06	1.09E-06	1.09E-06
10	5.26E-03	5.25E-03	-1.88E-06	-1.88E-06	1.23E-06	1.23E-06

The epot and ekin values are **identical** across all steps. The econs differences (last display digit) arise from different domain decomposition with 8 vs. 12 CPUs. With 32 CPUs the grid reduction would be $\sim 20\times$, bringing the previously OOM scenario (640 GB) well within a 560 GB memory budget.

11.8 Poisson MG Fine-Level Optimization

The multigrid Poisson solver (`poisson-mg`) is typically the single largest time consumer, accounting for 29–55% of total runtime. Several optimizations were applied to the fine-level V-cycle in `poisson/multigrid_fine.kjhan.f90`:

11.8.1 Precomputed Neighbor Grid Array

Before entering the V-cycle iteration loop, all 6-directional neighbor grids are precomputed and stored in a contiguous array:

```
! nbor_grid_fine(0:twondim, 1:ngrid)
!   index 0 = self (igridd_amr)
!   index 1..6 = neighbor grids in -x,+x,-y,+y,-z,+z
call precompute_nbor_grid_fine(ilevel)
```

This replaces per-cell `morton_nbor_grid` calls inside the Gauss-Seidel and residual loops with simple array lookups, eliminating hash table probes from the innermost loops.

11.8.2 Merged Red-Black Gauss-Seidel Exchange

The standard multigrid implementation performs a ghost zone exchange after each half-sweep of the red-black Gauss-Seidel smoother:

`red → exchange → black → exchange`

This was simplified to:

`red → black → exchange`

The black sweep uses slightly stale ghost values from the red sweep (“chaotic relaxation”), which does not affect multigrid convergence. This reduces the number of MPI exchanges per iteration from 9 to 5 (a 44% reduction in communication calls).

11.8.3 Residual and Norm Single Pass

The residual computation and the L^2 norm reduction were fused into a single pass:

```
! Optional norm2 argument: if present, compute
! both residual and L2 norm in one sweep
call cmp_residual_mg_fine(ilevel, norm2)
```

This eliminates a redundant loop over all cells when both the residual and its norm are needed (at the first iteration and after post-smoothing).

11.8.4 Division to Multiplication

In the Gauss-Seidel fast path, the division by `dtwondim` was replaced with multiplication by a precomputed reciprocal:

```
real(dp) :: oneoverdtwondim
oneoverdtwondim = 1.0d0 / dble(twondim)
! In loop:
phi = sum_neighbors * oneoverdtwondim ! was: / dtwondim
```

11.8.5 Performance Impact

These optimizations combined reduce the Poisson solver's share of total runtime:

Metric	Value
Before optimization	55.1% of runtime
After optimization	38.6% of runtime
MPI exchange reduction	44% fewer calls
Iteration count	unchanged (Level 8: 5, Level 9: 4)
Convergence	verified (econs = 3.79E-03 at step 10)

Chaotic Relaxation

The merged red-black exchange introduces a minor change in the energy conservation value ($3.77\text{E-}03 \rightarrow 3.79\text{E-}03$) due to the slightly different relaxation path. This is well within the multigrid tolerance and does not affect physical results. The membal and nomembal tests produce identical values.

12

Scaling Performance

Strong scaling tests were performed using a 200×10^6 particle cosmological simulation (levemin=8, levelmax=10), restarted from an HDF5 checkpoint at coarse step 5 and evolved to step 10 (5 coarse steps). The test platform is a dual-socket AMD EPYC 7543 node (64 physical cores, 128 threads via SMT) with 1 TB of DDR4 memory. The code was compiled with Intel MPI (mpiifx) and OpenMP (-qopenmp).

The variable-ncpu restart allows the 12-rank checkpoint to be read with any number of MPI ranks. A forced `load_balance` on the first coarse step ensures optimal grid distribution.

12.1 Pure MPI Scaling

All runs use `OMP_NUM_THREADS=1`. Speedup is relative to the 2-rank baseline.

Ranks	Elapsed	Speedup	Coarse	Particle	Poisson	MG	Hydro-GZ	Godunov	LoadBal
	(s)		(s)	(s)	(s)	(s)	(s)	(s)	(s)
2	193.0	1.00×	2.27	20.92	18.22	102.69	0.52	46.67	9.14
4	154.8	1.25×	2.59	16.74	14.76	82.00	0.89	36.64	7.62
8	88.9	2.17×	2.12	8.93	8.18	43.46	1.09	21.68	4.90
16	58.2	3.31×	2.56	4.67	4.97	23.90	1.26	15.88	3.96
32	34.7	5.57×	2.04	2.27	2.73	12.59	1.11	8.71	2.27
64	25.9	7.45×	1.91	1.18	1.64	7.95	1.13	5.14	1.61

12.2 Hybrid MPI + OpenMP Scaling

All configurations use a fixed total of 64 physical cores. Speedup is relative to the 2-rank pure MPI baseline (193.0 s).

Ranks×Thr	Elapsed	Speedup	Coarse	Particle	Poisson	MG	Hydro-GZ	Godunov	LoadBal
	(s)		(s)	(s)	(s)	(s)	(s)	(s)	(s)
64×1	25.9	7.45×	1.91	1.18	1.64	7.95	1.13	5.14	1.61
32×2	22.2	8.71×	1.55	1.37	1.64	7.23	0.82	5.19	1.62
16×4	24.9	7.76×	1.49	2.23	2.31	8.81	0.66	5.22	2.40
8×8	30.2	6.40×	1.26	3.80	3.34	11.35	0.68	5.46	3.00
4×16	44.4	4.34×	1.35	6.99	5.66	17.22	0.67	6.21	5.22

12.3 Key Observations

- **Optimal configuration** is 32 ranks \times 2 threads (22.2 s), which is 14% faster than 64 pure MPI ranks (25.9 s) and 8.7 \times faster than the 2-rank baseline.
- **Multigrid Poisson solver (MG)** dominates runtime and scales well: 102.7 s (2 ranks) \rightarrow 7.2 s (32 \times 2), a 14.2 \times speedup.
- **Godunov hydro solver** scales effectively with both MPI and OpenMP: 46.7 s (2 ranks) \rightarrow 5.2 s (32 \times 2), a 9.0 \times speedup.
- **Ghost-zone exchange (Hydro-GZ)** increases slightly with MPI rank count (0.52 s at 2 ranks \rightarrow 1.13 s at 64 ranks) due to more MPI messages, but *decreases* in hybrid mode (0.82 s at 32 \times 2) since fewer ranks means fewer messages.
- **Hybrid advantage:** 32 \times 2 beats 64 \times 1 because halving the MPI rank count reduces communication overhead (ghost zones, load balancing) while the 2 OpenMP threads provide sufficient intra-rank parallelism for compute-bound kernels.
- **Diminishing returns from OpenMP:** beyond 4 threads per rank, OpenMP thread synchronization overhead outweighs the communication savings, causing performance to degrade (4 \times 16: 44.4 s, worse than 8 \times 1: 88.9 s with only 8 total cores).
- **Load balancing** scales well: 9.1 s at 2 ranks \rightarrow 1.6 s at 64 ranks. This is a per-nremap cost (default nremap=5).
- **Physics correctness:** all configurations produce identical $e_{\text{pot}} = -1.88 \times 10^{-6}$ and $e_{\text{kin}} = 1.23 \times 10^{-6}$ at step 10.

Recommended Configuration

For a single dual-socket AMD EPYC node with 64 cores, we recommend **32 MPI ranks \times 2 OpenMP threads** for cuRAMSES cosmological simulations with $\sim 200 \times 10^6$ particles. Set OMP_NUM_THREADS=2, OMP_PROC_BIND=close, OMP_PLACES=cores, and I_MPI_PIN_DOMAIN=omp in the job script.

12.4 Per-Routine Scaling Analysis

Table 12.1 shows the complete timer breakdown for pure MPI scaling, and Table 12.2 for the hybrid configurations. All times are averages across ranks in seconds.

12.4.1 Scaling Categories

The routines fall into three distinct categories by their scaling behavior:

1. **Compute-bound (excellent MPI scaling).** courant (18.6 \times), particles (17.7 \times), flag (14.7 \times), poisson-mg (12.9 \times), poisson (11.1 \times), godunov (9.1 \times). These routines perform per-cell or per-particle work and scale nearly linearly with MPI rank count. In the hybrid regime, they are insensitive to the MPI/OMP ratio as long as total cores ≥ 32 .
2. **Communication-bound (anti-scaling).** hydro-gz (ghost zone exchange) grows from 0.52 s to 1.13 s as ranks increase from 2 to 64, because more ranks means more boundary surfaces and MPI messages. In the hybrid regime, reducing ranks to 32 or fewer recovers some of this cost (0.82 s at 32 \times 2, 0.67 s at 4 \times 16). io (HDF5 output) degrades from 3.14 s at 2 ranks to 6.17 s at 64 ranks due to MPI collective I/O contention. This is the clearest anti-scaling component: halving the rank count from 64 to 32 nearly halves the I/O time (6.17 s \rightarrow 3.37 s).
3. **Flat (Amdahl-limited).** coarse levels remains at ~ 2 s regardless of configuration — this work is inherently serial or involves global communication at the coarsest level. It

Table 12.1: Full timer breakdown — pure MPI (`OMP_NUM_THREADS=1`).

Routine	2r	4r	8r	16r	32r	64r	Scale
poisson-mg	102.69	82.00	43.46	23.90	12.59	7.95	12.9×
godunov	46.67	36.64	21.68	15.88	8.71	5.14	9.1×
particles	20.92	16.74	8.93	4.67	2.27	1.18	17.7×
poisson	18.22	14.76	8.18	4.97	2.73	1.64	11.1×
flag	11.35	8.23	4.41	2.78	1.28	0.77	14.7×
loadbalance	9.14	7.62	4.90	3.96	2.27	1.61	5.7×
io	3.14	1.96	1.86	2.72	4.23	6.17	0.5×↓
coarse levels	2.27	2.59	2.12	2.56	2.04	1.91	1.2×
courant	2.23	1.49	0.76	0.41	0.23	0.12	18.6×
hydro-gz	0.52	0.89	1.09	1.26	1.11	1.13	0.5×↓
rev-gz	0.53	1.79	2.42	1.75	0.78	0.71	0.7×↓
set unew	0.28	0.26	0.22	0.23	0.15	0.10	2.8×
set uold	1.11	0.77	0.38	0.24	0.15	0.10	11.1×
refine	0.61	0.86	0.69	0.62	0.43	0.27	2.3×
TOTAL	219.7	176.6	101.1	66.0	39.0	28.8	7.6×

Table 12.2: Full timer breakdown — hybrid MPI+OpenMP (fixed 64 cores).

Routine	64×1	32×2	16×4	8×8	4×16
poisson-mg	7.95	7.23	8.81	11.35	17.22
godunov	5.14	5.19	5.22	5.46	6.21
particles	1.18	1.37	2.23	3.80	6.99
poisson	1.64	1.64	2.31	3.34	5.66
io	6.17	3.37	2.25	2.02	1.89
loadbalance	1.61	1.63	2.46	3.06	5.27
flag	0.77	0.89	1.50	2.47	4.71
hydro-gz	1.13	0.82	0.66	0.68	0.67
rev-gz	0.71	0.57	0.83	0.56	0.58
coarse levels	1.91	1.55	1.49	1.26	1.35
courant	0.12	0.13	0.13	0.13	0.14
set uold	0.10	0.10	0.10	0.11	0.14
set unew	0.10	0.10	0.09	0.09	0.10
refine	0.27	0.28	0.31	0.25	0.40
TOTAL	28.8	24.9	28.4	34.6	51.3

becomes a larger fraction of runtime as other components scale (from 1% at 2 ranks to 6.6% at 64 ranks).

12.4.2 Bottleneck Shift

As the rank count increases, the dominant bottleneck shifts:

- At 2 ranks: `poisson-mg` (46.7%), `godunov` (21.2%), `particles` (9.5%)
- At 32×2 : `poisson-mg` (29.1%), `godunov` (20.9%), `io` (13.6%)
- At 64×1 : `poisson-mg` (27.6%), `io` (21.4%), `godunov` (17.8%)

At high rank counts, I/O emerges as the second-largest cost, suggesting that asynchronous or node-local I/O strategies would yield further speedup.

13

GPU Hydro Acceleration

cuRAMSES supports GPU acceleration for the main hydro solver routines via CUDA. The implementation uses a **dynamic hybrid CPU/GPU dispatch** model where OpenMP threads compete for a shared pool of GPU streams: threads that acquire a stream offload their work to the GPU, while the remaining threads continue on the CPU. This design naturally adapts to any ratio of CPU cores to GPU streams.

13.1 Architecture Overview

The GPU acceleration is enabled at build time with `make HDF5=1 USE_CUDA=1`. At runtime, each MPI rank initializes a pool of 8 CUDA streams, distributed across available GPUs by local rank (`local_rank % device_count`). The multi-architecture binary supports sm_80 (A100), sm_86 (A10), and sm_90 (H100) targets simultaneously.

Key Files

- `cuda_stream_pool.{h,cu}` — Thread-safe stream pool with atomic acquire/release
- `hydro_cuda_kernels.cu` — CUDA kernels: `ctoprim`, `uslope`, `trace3d`, `flux`, `difmag`, `synchro`, `cmpdt`, `gradient_phi`, `upload_fine`
- `hydro_cuda_interface.f90` — Fortran ISO_C_BINDING wrappers with assumed-size buf(*) for safe `c_loc` usage
- `cuda_commons.f90` — Fortran module exposing stream pool to the AMR code

13.2 Hybrid CPU/GPU Dispatch

The dispatch model is illustrated for the Godunov solver (`godunov_fine.kjhan.f90`), which accounts for ~30% of total runtime:

Listing 13.1: Dynamic hybrid dispatch pattern

```
1 !$omp parallel
2     stream_slot = cuda_acquire_stream()    ! atomic, -1 if busy
3 !$omp do schedule(dynamic)
4     do igridd = 1, ncache, nvector
5         if (stream_slot >= 0) then
6             ! GPU: gather into superbatches buffer
7             ! Flush when buffer full -> async GPU kernel
8         else
9             ! CPU: standard godfine1 computation
10        end if
11    end do
```

```

12 !$omp end do nowait
13 ! GPU thread: flush remaining buffer, release stream
14 !$omp end parallel

```

The `schedule(dynamic)` clause ensures load balancing: if one thread is waiting for GPU synchronization, other threads pick up the remaining iterations on CPU. This is particularly effective when the number of OMP threads exceeds the number of GPU streams.

13.2.1 Superbatch Buffering

GPU kernel launches have significant latency ($\sim 10\text{--}50 \mu\text{s}$). To amortize this, each GPU thread accumulates grid data into a **superbatch buffer** of size `SUPER_SIZE` (typically 4096 grids) before launching a single kernel covering all accumulated grids. The buffer is flushed when it approaches capacity or when the `omp do` loop completes.

For the Godunov solver, the superbatch contains the full stencil data (conservative variables and gravitational acceleration on a $6 \times 6 \times 6$ grid per AMR oct), and the GPU pipeline executes 5 kernels in sequence: `ctoprim` → `uslope` → `trace3d` → `flux` → `difmag`.

13.2.2 Accelerated Routines

Six routines use the hybrid dispatch pattern:

Routine	Kernel	Data per cell
<code>godunov_fine</code>	5-kernel pipeline	Full stencil ($6^3 \times \text{nvar}$)
<code>synchro_hydro_fine</code>	Gravity kick	$\rho, \mathbf{p}, E, \mathbf{f}$ (8 doubles)
<code>courant_fine</code>	CFL timestep	<code>nvar + ndim</code> (14 doubles)
<code>force_fine</code>	Gradient of ϕ	$3^3 \times \phi + 3^3 \times \mathbf{f}$
<code>upload_fine</code>	Prolongation	Child + parent cells
<code>cooling_fine</code>	Cooling function	Temperature + density

13.2.3 Memory Management

Each stream slot has persistent device memory that is lazily allocated and grown with 2× over-allocation:

- **Hydro buffers:** `d_uloc`, `d_gloc`, `d_flux`, `d_tmp`, `d_ok` — sized for `SUPER_SIZE` grids
- **Generic buffers:** Per-stream device memory for lightweight kernels (`synchro`, `cmpdt`, etc.)

The host-side superbatch buffers are allocated per-thread via OMP private clauses and passed to external flush subroutines as explicit-shape arrays to avoid Intel ifx compiler issues with `c_loc` on assumed-shape arrays.

13.3 Fortran–CUDA Interface

A critical design choice is the two-layer interface between Fortran and CUDA:

1. **C binding layer** (`_c` suffix): Direct `bind(C)` interface passing type(`c_ptr`) for array arguments.
2. **Fortran wrapper layer** (`_f` suffix): Receives assumed-size `buf(*)` arrays and converts to `c_ptr` via `c_loc(buf(1))`.

The wrapper layer is necessary because Intel ifx may return incorrect values from `c_loc` on assumed-shape arrays. The assumed-size `buf(*)` pattern avoids array descriptors entirely and guarantees that `c_loc` returns the correct data address.

13.4 Multi-GPU Support

When multiple GPUs are available on a node, each MPI rank is assigned to a GPU based on its local rank (`local_rank % device_count`). The stream pool calls `cudaSetDevice(g_device_id)` in `cuda_acquire_stream` to ensure each OMP thread operates on the correct GPU context.

The CUDA binary is compiled with multiple architecture targets:

```
1 CUDA_ARCH = -gencode arch=compute_80,code=sm_80 \
2                 -gencode arch=compute_86,code=sm_86 \
3                 -gencode arch=compute_90,code=sm_90
```

This enables a single binary to run on A100 (sm_80), A10 (sm_86), and H100 (sm_90) GPUs without recompilation.

13.5 Verification

The GPU-accelerated code produces bit-identical physics results compared to the CPU-only build:

- 4 ranks \times 1 GPU (A10): $e_{\text{cons}} = 5.23 \times 10^{-3}$, $e_{\text{pot}} = -1.88 \times 10^{-6}$, $e_{\text{kin}} = 1.23 \times 10^{-6}$ at step 10
- 4 ranks \times 4 GPUs (A100): identical physics results with correct per-rank GPU assignment



Namelist Reference

14

RUN_PARAMS

Runtime control parameters.

Parameter	Type	Default	Description
cosmo	logical	.false.	Enable cosmological simulation
pic	logical	.false.	Enable Particle-In-Cell
poisson	logical	.false.	Enable Poisson gravity solver
hydro	logical	.false.	Enable hydrodynamics
rt	logical	.false.	Enable radiative transfer
sink	logical	.false.	Enable sink particles
verbose	logical	.false.	Verbose output
debug	logical	.false.	Debug mode
nrestart	integer	0	Restart file number (0 = new run)
nstepmax	integer	1000000	Maximum number of coarse steps
ncontrol	integer	1	Frequency of control variable output
nsubcycle	integer[]	2	Subcycling factor per level
nremap	integer	5	Load balancing frequency (0=never)
ordering	char	hilbert	Domain decomposition: hilbert, bisection, ksection
static	logical	.false.	Static (no refinement) mode
overload	integer	1	MPI overload factor
cost_weighting	logical	.true.	CPU time-based cost weighting
<i>Memory-based load balancing (new)</i>			
memory_balance	logical	.false.	Enable memory-weighted balancing
mem_weight_grid	integer	270	Memory per grid (dp-equivalents)
mem_weight_part	integer	12	Memory per particle (dp-equivalents)
<i>Job control (new)</i>			
jobcontrolfile	char(128)	''	Runtime control file for stop/output requests

Job control file. When `jobcontrolfile` is set to a non-empty path, rank 0 reads the file at every coarse step. Each line contains two integers: `step_number` and `action_code`.

- `step_number` = 0: match at every step (immediate).
- `step_number` = N : match when `nstep_coarse` = N .
- `action` = 1: write an extra output and continue.
- `action` = -1: write an extra output and stop gracefully.

Example: to request a graceful stop at the next coarse step, create the file with `echo "0 -1" > jobcontrol.txt`. The file is not deleted after reading; lines with `step_number = 0` re-trigger every step until the file is removed or modified.

15

AMR_PARAMS

Adaptive Mesh Refinement grid parameters.

Parameter	Type	Default	Description
levelmin	integer	1	Minimum (uniform) refinement level
levelmax	integer	1	Maximum refinement level
ngridmax	integer8	0	Max grids per CPU (0 = auto)
ngridtot	integer8	0	Total grids for auto-computation
npartmax	integer	0	Max particles per CPU (0 = auto)
nparttot	integer	0	Total particles for auto-computation
nexpand	integer[]	1	Mesh expansion layers per level
boxlen	real(dp)	1.0	Box side length in code units

16

OUTPUT_PARAMS

Output control parameters.

Parameter	Type	Default	Description
noutput	integer	1	Number of scheduled outputs
foutput	integer	1000000	Output every N steps
fbackup	integer	1000000	Backup every N steps
aout	real[]	1.1	Output expansion factors (cosmo)
tout	real[]	0.0	Output times (non-cosmo)
outformat	char(10)	original	Output format: original (per-CPU binary) or hdf5 (single HDF5 file)
informat	char(10)	original	Restart format: original or hdf5. Can differ from outformat
output_mode	integer	0	Hi-res output mode
gadget_output	logical	.false.	Write Gadget-format snapshots
walltime_hrs	real(dp)	-1	Job walltime (hours, < 0 = ignore)
minutes_dump	real(dp)	1.0	Dump this many minutes before walltime

17

INIT_PARAMS

Initial conditions parameters.

Parameter	Type	Default	Description
filetype	char(20)	ascii	IC file format: ascii, grafic, gadget
initfile	char[]	' '	IC file path per level
multiple	logical	.false.	Multiple IC files per rank
nregion	integer	0	Number of IC regions

18

REFINE_PARAMS

Refinement criteria parameters.

Parameter	Type	Default	Description
m_refine	real[]	-1	Lagrangian mass threshold per level
ivar_refine	integer	-1	Variable index for gradient refinement
var_cut_refine	real(dp)	-1	Variable threshold
mass_cut_refine	real(dp)	-1	Particle mass threshold
interpol_var	integer	0	Interpolation variable (0=conservative, 1=primitive)
interpol_type	integer	1	Interpolation type (0=MinMod, 1=MonCen)
sink_refine	logical	.false.	Fully refine around sinks
jeans_ncells	real(dp)	-1	Jeans length in cells (> 0 enables polytropic EOS)

19

HYDRO_PARAMS

Hydrodynamics solver parameters.

Parameter	Type	Default	Description
gamma	real(dp)	$\frac{5}{3}$	Adiabatic index γ
courant_factor	real(dp)	0.8	Courant–Friedrichs–Lowy number
scheme	char(20)	muscl	Hydro scheme (muscl)
slope_type	integer	1	Slope limiter (1=MinMod, 2=MonCen, 3=unlimited)
pressure_fix	logical	.false.	Pressure floor for strong shocks
beta_fix	real(dp)	0.0	Pressure fix strength
isothermal	logical	.false.	Isothermal mode

20 POISSON_PARAMS

Gravity and Poisson solver parameters.

Parameter	Type	Default	Description
epsilon	real(dp)	10^{-4}	Multigrid convergence criterion
gravity_type	integer	0	Gravity type (0=self-gravity, > 0 =analytic)
cg_levelmin	integer	999	Min level for CG solver fallback
cic_levelmax	integer	0	Max level for CIC particle interpolation

21 PHYSICS_PARAMS

Subgrid physics parameters (star formation, feedback, cooling).

Parameter	Type	Default	Description
cooling	logical	.false.	Enable radiative cooling
metal	logical	.false.	Enable metal tracking
haardt_madau	logical	.false.	UV background
z_reion	real(dp)	8.5	Reionization redshift
<i>Star formation</i>			
n_star	real(dp)	0.1	SF density threshold (H/cc)
t_star	real(dp)	0.0	SF timescale (Gyr)
eps_star	real(dp)	0.0	SF efficiency
T2_star	real(dp)	0.0	ISM polytropic temperature
g_star	real(dp)	1.6	ISM polytropic index
sf_birth_properties	logical	.true.	Output stellar birth properties
<i>Cosmology (read from IC header or namelist)</i>			
omega_b	real(dp)	0.0	Baryon density Ω_b
omega_m	real(dp)	1.0	Matter density Ω_m
omega_l	real(dp)	0.0	Dark energy Ω_Λ
h0	real(dp)	1.0	Hubble constant H_0 (km/s/Mpc)
<i>Feedback</i>			
f_ek	real(dp)	1.0	SN kinetic energy fraction
rbubble	real(dp)	0.0	SN superbubble radius (pc)
yieldtable- filename	char	—	Yield table file path

22

Example Namelist

22.1 Cosmological Simulation with Memory Balancing

`cosmo_ksection_membal.nml`

```
&RUN_PARAMS
cosmo=.true.
pic=.true.
poisson=.true.
hydro=.true.
nrestart=0
nremap=5
nsubcycle=1,1,2
ncontrol=1
nstepmax=10
ordering='ksection'
memory_balance=.true.
jobcontrolfile='jobcontrol.txt',
/

&OUTPUT_PARAMS
noutput=1
aout=1.0
/

&INIT_PARAMS
filetype='grafic'
initfile(1)='/path/to/ics/level_008',
/

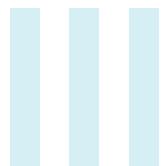
&AMR_PARAMS
levelmin=8
levelmax=10
nexpand=1
ngridtot=40000000
nparttot=200000000
/

&REFINE_PARAMS
m_refine=3*8.,
ivar_refine=0
interpol_var=1
interpol_type=0
/
```

```
&HYDRO_PARAMS
gamma=1.6666667
courant_factor=0.8
scheme='muscl'
slope_type=1
/

&POISSON_PARAMS
/

&PHYSICS_PARAMS
sf_birth_properties=.false.
yieldtablefilename='yield_table.asc'
/
```



Build and Testing

23

Build Instructions

23.1 Prerequisites

- Intel Fortran Compiler (`ifx`) with MPI support (`mpiifx`)
- OpenMP support (`-qopenmp`)

23.2 Build

Build Commands

```
cd bin  
make clean  
make
```

The binary is produced as `bin/ramses_final3d`.

23.3 Compile-Time Options

Key preprocessor flags set in the Makefile:

Flag	Meaning
<code>-DNDIM=3</code>	Three-dimensional simulation
<code>-DNVECTOR=32</code>	Vector length for grid sweeps
<code>-DLONGINT</code>	64-bit integer grid indices
<code>-DQUADHILBERT</code>	Quad-precision Hilbert keys
<code>-DNVAR=11</code>	Number of hydro variables
<code>-DNPRE=8</code>	Passive scalar variables

24

Verification Tests

24.1 Test Configuration

Setting	Value
IC	MUSIC level_008 (GRAFIC2 format)
MPI ranks	12
Levels	8–10
Steps	10
Ordering	ksection

24.2 Reference Values (nremap=5)

Test	nparttot	econs	epot	ekin	mcons
membal	200M	3.77E-03	-1.88E-06	1.23E-06	-1.84E-16
nomembal	80M	3.77E-03	-1.88E-06	1.23E-06	-1.84E-16

24.3 Running Tests

Membal Test

```
cd test_ksection/run_cosmo_membal
cp ../../bin/ramses_final3d .
mpirun -np 12 ./ramses_final3d \
    ./cosmo_ksection_membal.nml
```

Verify that at step 10: econs=3.77E-03, epot=-1.88E-06, ekin=1.23E-06.



Modified Files Summary

File	Modifications
<i>patch/cuda/</i>	
amr_parameters.jaehyun.f90	Memory balance params, nremap=5 default
amr_commons.kjhan.f90	grid_level array, communicator type
read_params.jaehyun.f90	Read new namelist params
bisection.f90	nc_in parameter, 64-bit histograms
ksection.f90	K-section tree + exchange routines
load_balance.kjhan.f90	numbp sync, bulk exchange, internal timing
virtual_boundaries.kjhan.f90	Ksec ghost exchange + bulk + build_comm
multigrid_fine_commons.f90	MG ksection communication
init_amr.f90	Memory savings, Morton init
update_time.f90	Memory reporting
adaptive_loop.jaehyun.f90	writemem_minmax, Morton rebuild, bulk exchange
<i>patch/oct_tree/</i>	
morton_keys.f90	Morton key computation
morton_hash.f90	Hash table + neighbor helpers
morton_init.f90	Build and verify
refine_utils.f90	Hash maintenance in grid ops
nbors_utils.kjhan.f90	Morton-based neighbor lookup
<i>patch/Horizon5-master-2/</i>	
init_flow_fine.f90	Stream access IC reading
init_part.f90	Stream access particle IC
particle_tree.kjhan.f90	MPI_ALLTOALL → ksection
amr_step.jaehyun.f90	Bulk virtual exchange
feedback.kjhan3.f90	SNII feedback spatial binning
<i>patch/cuda/ (HDF5 I/O — new files)</i>	
ramses_hdf5_io.f90	HDF5 wrapper module (parallel create/open/write/read)
backup_hdf5.f90	HDF5 output: AMR, hydro, gravity, particles, sinks
restore_hdf5.f90	HDF5 restart: AMR tree rebuild + data restore
output_amr.kjhan.f90	HDF5 dispatch in dump_all